FedCast

A forecaster for the next Federal Reserve rate decision, built to be judged by its process rather than its luck

Tan Kah Ming · Fullerton · October 2026

About me



Tan Kah Ming

Senior AI, Data and Technology Consultant, seconded to Sovereign Wealth Fund.

Before that

- Data scientist, RSAF Data Office: contributed to RSAF AI Strategy, set up its first digital unit
- Air engineering and weapons logistics officer, where the data-driven habit started (Weibull analysis, predictive maintenance)

Training

M.Eng Aeronautical Engineering, Imperial College London · Graduate Certificate in Predictive Analytics, NUS · CFA Level 3 candidate (February 2027)

Things I have built that models/Applications in

1

Investment analytics, valuation and execution

Investment analytics: AI-assisted performance attribution and idea tracking for an investment team.
Valuation: Machine-learning valuation for illiquid assets.
Execution: Trade workflow automation for investment teams.

2

Data Science

- 4th of 160+ teams (MIT, DoD, USAF) in the 2023 USAF AI Accelerator, predicting pilot flying difficulty from physiology;
- An ammunition recommender (quantile LightGBM) that cut wastage with no shortfall -> Defence Manpower Committee Innovation Award 2022.

3

COVID-19, 2021

Built, in weeks, the system that engaged every home-recovery patient in Singapore, 10,000 people a day, on form.sg, Azure and Zendesk. The lesson that carries over: a system people rely on must say what it knows and what it does not.

The question, the answer, and what happened

28 Oct

the decision being forecast, 2026

67.6%

FedCast, no change: ledger entry #18, nine hours before

Hold

what the Fed did: 3.75 to 4.00%, unchanged

10 to 2

two dissents, both for the +25 bps the specialists leaned toward

How the number was made

- One frozen snapshot of 23 evidence items, pulled 28 October: futures, eight FRED series, both statements, minutes, transcript, projections, speeches, 32 years of decisions, the NY Fed natural rate, prediction-market odds
- Four models pooled by code at fixed weights (market 0.6, Taylor rules 0.2, odds 0.1, base rates 0.1), then a 1.5% floor on every outcome: no change 70.9%, hike 26.2%
- Five specialist analysts, every claim quoted and verified by code, lean the number +4.0 bp toward a hike: no change 67.6%, +25 bps 22.4%, +50 bps 7.7%
- Hash-chained ledger entry #18; recomputes exactly from its snapshot and the analyst runs it names

What the outcome says

- The most likely outcome was the outcome. The market alone was more confident (90/10); at five weeks we do not claim to beat it
- What the system added was the reasons: the hawkish lean came from the communications, minutes and inflation readers; the labour reader was the lone dove. Two dissents for a hike is that lean, in the vote
- Every reason is a quote code found in the frozen evidence, so the call can be taken apart claim by claim, which the rest of this talk does
- One decision proves little. It is scored formally (log score, Brier, against the market and two naive baselines) the moment the NY Fed's post-meeting rate row lands, and every weight decision waits on that record

Live demo 1: the forecast

1 min

Open

- The app at localhost:8765, Forecast view
- Top bar: when the evidence was pulled, what is new at the sources

Point at

- The two tracks side by side: machine only, and with views
- The meta line: ledger entry, snapshot hash, the specialists' tilt
- The five bars

Say

- This number was made at <time> from evidence frozen at <time>. Everything I show next is how it got there.

LIVE DEMO

Four kinds of signal, and where FedCast sits

Market signals

Fed funds futures (CME FedWatch); prediction markets (Polymarket, Kalshi)

What it is Prices and crowd odds on each outcome

Strength Futures about 88% right at 30 days, 1994 to 2024; odds move fast

Economic signals

Economist surveys; the macro data itself

What it is The median of polled forecasters; the data the committee reads

Strength Broad human judgement; the data is the Fed's own input

Policy rules

The Taylor family

What it is The rate a formula says is appropriate given inflation and unemployment

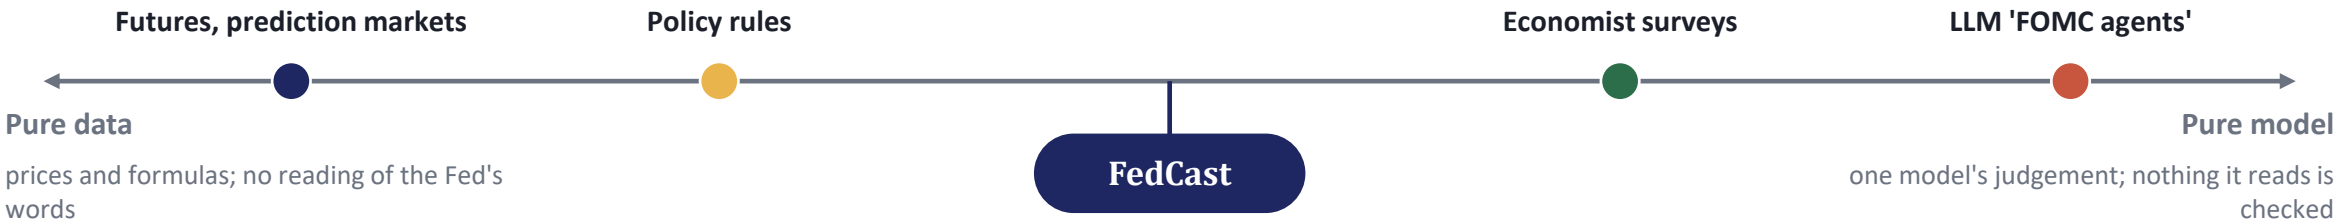
Strength Transparent; the Fed's own published benchmarks

ML / AI signals

LLM 'FOMC agent' papers (2025)

What it is Models that read the Fed's text and vote

Strength Reported 90%+ accuracy on 2023 to 2024



Takes one input from each area, prices, 32 years of decisions, the Fed's own rules and a language model that must quote what it reads, and combines them by code at fixed weights. Neither a formula nor an oracle.

Sources: Bonini, Huang and Simaan, 'Watching the FedWatch' (232 meetings, May 1994 to March 2024); FedSight AI preprint, December 2025; Takano, Hirano and Nakagawa, arXiv 2511.02469.

Why this approach ?

1 We judge the process, not the luck

- Eight decisions a year: a good year proves nothing
- The yardstick is a ten-point scorecard, fixed before any code: same evidence, same answer? every claim quoted? any number from memory? a shock moves the number the right way?
- Outcomes are logged after each decision and never used to tune

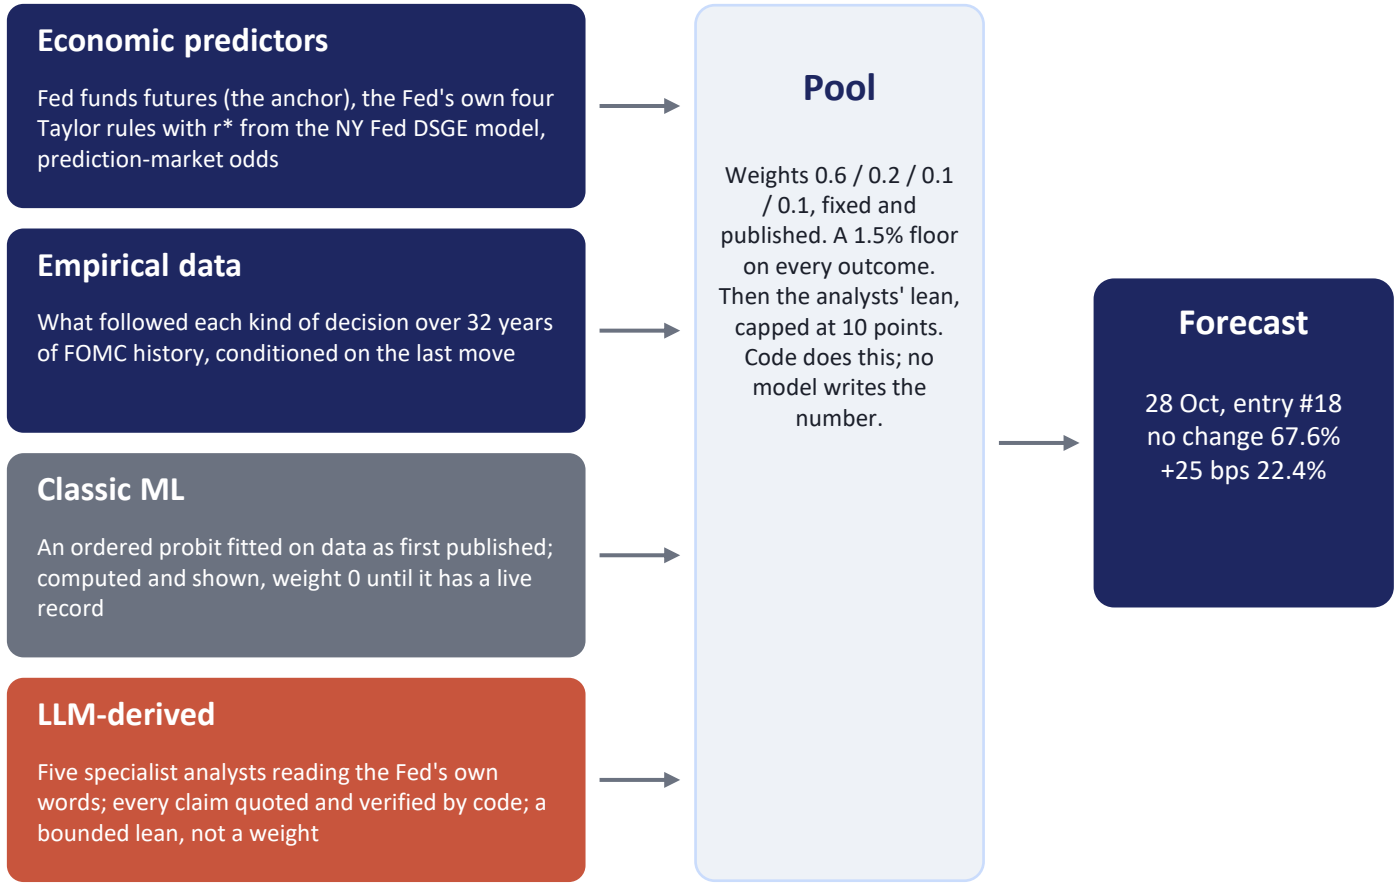
2 Every step is auditable

- An analyst that says 'inflation remains elevated' must quote the statement line it read, and code checks the quote is there, character for character
- Every number it uses must exist in the frozen evidence, or it is flagged as from memory
- Every model lists the items it read; every ledger entry recomputes from its snapshot. If it is not quoted, it does not count

3 A harness that compounds

- Everything is quotable and the analysts can read only the evidence, so a wrong call is a traceable wrong call: we see which quote led where
- The harness is repeatable, so it goes wrong the same way every time; a fix can be tested, and the drift table shows it stayed fixed
- That is the virtuous cycle: audit, fix, re-measure. The ladder, the scorecard and the drift table later in the deck are that cycle

Model Architecture



The chain, in order

Evidence — 23 items frozen and fingerprinted, dated in US time, data as first published. Nothing outside it may be read.

Signals — Computed by code and handed to the analysts as numbers: inflation gap, labour slack, financial conditions, statement tone, base rates, projections.

Pool — Four families at fixed weights, the floor, then the specialists' lean.

Analysts — Only analysts 4 to 8 enter the number. Analysts 1 to 3 are checks and balances on the ladder.

You — Views tilt by at most 10 points and expire; watch-outs must be answered by id; documents and odds become evidence.

Ledger — Append-only, hash-chained; every entry recomputes from its snapshot and the runs it names.

The arithmetic: from five models to one number

Model	What it reads	−25	no change	+25	+50	weight
Market-implied	November fed-funds future against the overnight rate: +2.5 bp expected, split between the two nearest outcomes	0	90	10	0	0.6
Taylor family	12-month core PCE 3.0%, unemployment 4.2%, r* from the NY Fed DSGE model (0.65 / 1.86 / 3.06%); three level rules × nine r*/u* pairs plus the first-difference rule × three u*: 30 prescriptions, mean +22.5 bp	4.5	26.9	44.6	23.8	0.2
Prediction market	Polymarket, 26 Oct, $67m traded: 94.5% no change at the mid price, 94.9% once the five prices are renormalised	0.2	94.9	4.5	0.4	0.1
Historical base rates	What followed a hike at the 49 scheduled meetings since 1994 that came after one	0	34.7	52.2	13.1	0.1
Ordered probit v2	Nine features on first-published data; failed v1, caveated v2	0	24.2	63.8	12.0	0
Pooled	Weighted average, then the 1.5% floor	**1.5**	**70.9**	**20.2**	**6.0**	
After pooling	**Specialists' lean: a modifier on the pooled number, not a model in the pool.** Mean of analysts 4 to 8, +4.0 bp hawkish = 4.0 points of probability moved toward a hike, capped at 10.	−0.3	−3.3	+2.2	+1.7	—
Machine forecast	Ledger entry #18, 28 October	**1.2**	**67.6**	**22.4**	**7.7**	

Weights are based on judgement

Logged as amendments, never fitted. Ideally we would backtest the appropriate weights; with eight decisions a year there is no record to fit on yet, so live scores from 28 October onward decide them, by amendment.

Live demo 2: the evidence graph and the pool

3 min

Open

- Evidence graph
- Click the Market-implied box
- Click the Pool box
- Methodology, scroll to the worked example

Point at

- Market-implied: the futures price, 100 minus price, the expected change, the 90/10 split
- Pool: the four weights, the pooled line, the specialists' tilt line, Machine track
- Any evidence item: the frozen text and the link to the original

Say

- Blue is in the number, grey is computed but not in it. Click a model and you see only what it read.
- The language model's lean is the one line between Pooled and Machine track. It can tilt, it cannot write.

LIVE DEMO

Working around the model's limits

Three things a language model cannot be made to stop doing. None is removed; each is designed around, and the design is measured on every snapshot.

LIMIT

Memory cannot be switched off

Everything it knows about the Fed is in its weights, and it will use it unless something stops it.

WHAT WE DO

- **Prompt engineering:** use only the evidence; do not use anything you remember about the economy, the Fed or markets.
- **Answer shape:** every claim must carry a verbatim quote from a named item, so a remembered fact has nowhere to stand.
- **Harness:** every number in the answer is matched to the snapshot by code; one that is not found is flagged as from memory and sent back once to be removed.

PROOF, 28 OCTOBER

38 → 0 figures in the answer that appear nowhere in the evidence: analyst 1, which sees no evidence, against analysts 2 to 8, five runs each

LIMIT

The same question gets different answers

Ask twice and the number moves.

WHAT WE DO

- **Answer shape:** a JSON schema fixes the form of the answer; the runtime rejects anything else before anyone reads it.
- **Five identical runs** per analyst per snapshot: the same evidence, the same question, the same schema, only the run number differs. Each is stored and hashed, so it can be audited later.
- **Repeatability is measured:** how far the answer moves between those five runs. Pass mark: 2 points.

PROOF, FOUR SNAPSHOTS

4.2 → 0.8 how far the answer moves between identical runs: naive 4.2 points, specialists 0.8 bp, mean of four snapshots

LIMIT

It can invent its own “evidence”(hallucinate)

A fluent quote or figure that was never in the document.

WHAT WE DO

- **Harness:** every quote and every number in a claim is verified by code, not by a model. A quote must match its item character for character; a quote stitched from two places fails.
- **Bounds:** every watch-out must be answered by id; the direction call is clipped to ±10 bp.
- **One repair round** with the verifier's report; both answers are kept. What still fails is recorded as unverified, never hidden.

PROOF, 28 OCTOBER

326 / 327 claims whose quote code found verbatim, analysts 3 to 8, 30 runs; the one miss is on the record

LLM breakdown: the group of analysts

CONTROLS — BUILD THE LADDER

1 — Naive

The question and the date. Nothing else.

Intent: what the model does on its own, from memory. Its output also forms the baseline for every other step.

2 — Evidence shown, nothing verified

Every document and pre-read, 128,000 characters. Free-form answer.

Intent: to observe what evidence alone fixes (memory), and what it does not (variance; nothing checkable).

3 — Generalist with verifier

Same evidence; a fixed answer shape, a quote per claim, code checks every quote and number.

Intent: isolate the value of the structure and verification.

SPECIALISTS · THE FIVE THAT ENTER THE NUMBER — IN THE FORECAST

Analyst 3's shape and verifier, with one change: each reader gets **one slice of evidence**. A judgement is attributable to its slice, and a dissent survives instead of being averaged away.

4 — Communications

Both statements, the minutes, the transcript, the speeches, the tone pre-read.
Intent: what the committee is signalling.

5 — Minutes digest

The minutes only.
Intent: the committee's own account of its debate, read in isolation.

6 — Inflation

Projections table, both statements, the inflation pre-read.
Intent: is inflation pulling toward a hike.

7 — Labour market

Projections table, both statements, the labour pre-read.
Intent: is labour pulling the other way.

8 — Financial conditions

Both statements, the financial-conditions and base-rate pre-reads.
Intent: is the market or stress already doing the committee's work.

Each specialist is asked one question about its own slice (above), never for a probability. It returns claims, each with a verbatim quote and item id, and a lean of at most 10 bp either way: direction and strength, positive toward a hike. Code averages the five leans; each point of lean moves one point of probability, capped at 10.

Every analyst runs five times on the same snapshot. For analysts 1 to 3 those runs are not forecasts: they build the benchmark that 4 to 8 are measured against.

1 sets the baseline: variance of the probabilities across runs; numbers from memory (all of them); nothing quoted.

2 tests immutable evidence: numbers from memory fall to zero; the variance does not; nothing can be checked.

3 tests the answer shape and code verification: variance falls under 1 bp; every claim quoted and checked.

Live demo 3: what each analyst is asked and given

2 min

Open

- Analyst prompts view
- Analyst 1: open the full prompt
- Analyst 6: open the full prompt

Point at

- Analyst 1: 195 characters, no evidence
- Analyst 6: the system prompt, the enforced answer schema, the inflation pre-read, the item ids
- The prompt hash: five repeats share it

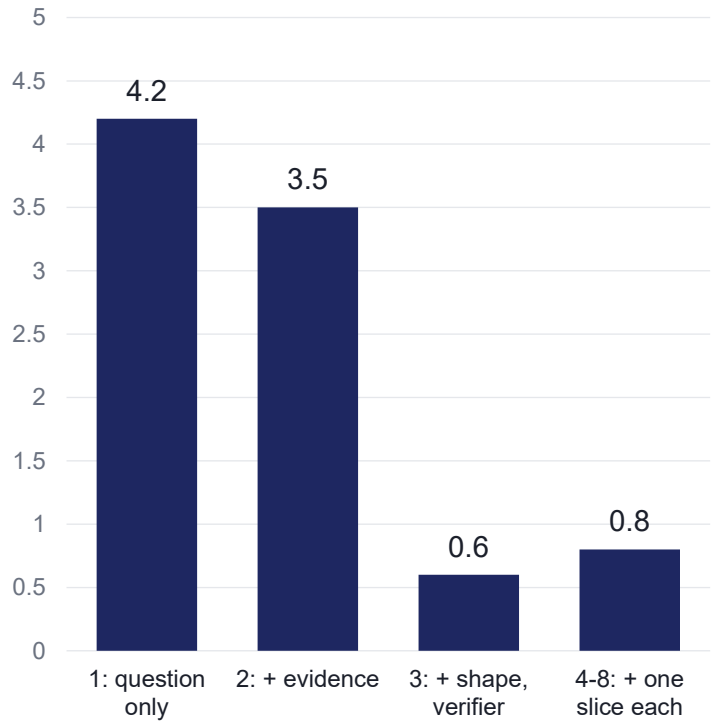
Say

- If it is not on this page, the analyst does not get it.
- Same model. The only difference between analyst 1 and analyst 6 is this page.

LIVE DEMO

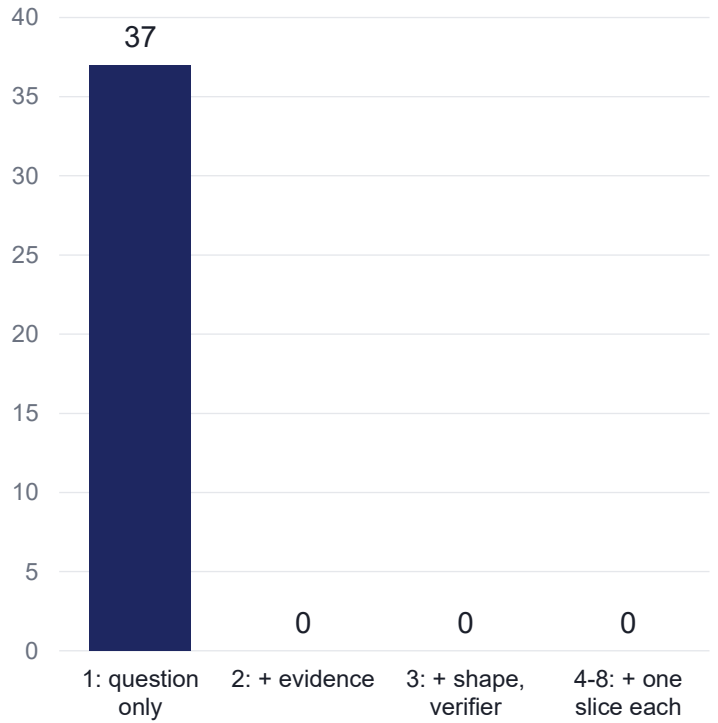
How good is the harness: what each step bought

Repeatability: variance between identical runs (pp)



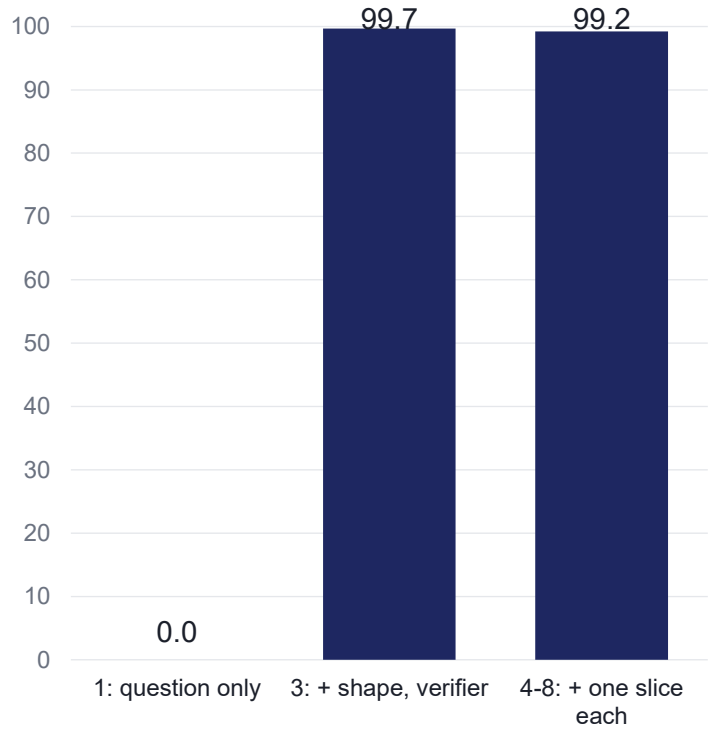
Largest standard deviation across five runs, mean of four snapshots. Pass mark 2 pp. Analysts 3 to 8: standard deviation of the lean, 1 bp = 1 pp.

Numbers from memory (count)



Numbers in the answer that exist nowhere in the frozen evidence, summed over five runs. Pass mark: exactly 0.

Claims with a verbatim quote (%)



Share of claims whose quote code found character for character in the item named. Pass mark 95%. Analyst 2 is not shown: it quotes nothing, so it cannot be measured.

The step from analyst 2 to analyst 3 is the harness: a fixed answer shape with a quote per claim, checked by code. Specialists keep both and add attribution. What they concluded, four snapshots: analyst 1 a cut at 57 to 67%; analyst 2 +25 bp at 50 to 57%; analyst 3 +5.8 to +7.0 bp hawkish; specialists +4.0 to +4.4 bp, labour the lone non-hawk.

LLM analysts: outcome

Each specialist runs five times on the same frozen snapshot. The five answers must sit within the variance pass mark (2 points), every claim must carry a quote that code finds in the evidence, character for character, and no number may come from memory. Only then does the lean count toward the number.

25 runs	25 prompts	263 claims	263 verified	5 leans, ±10 bp
5 specialists × 5 identical runs	one slice each, hashed	every one with a verbatim quote	100%; 0 numbers from memory	+6.0 · +6.2 · +5.8 · −1.4 · +3.4 bp

WHAT THE FIVE SAID, AND WHAT IT DID TO THE NUMBER

Communications	**+6.0 hike**
Minutes	**+6.2 hike**
Inflation	**+5.8 hike**
Financial conditions	**+3.4 hike**
Labour market	**−1.4 hold**
Equal-weight mean	**+4.0 bp**

Four say hike, one says hold back. Plain average, one vote each: +4.0. **Each point of lean moves one point of probability** from the hold side to the hike side, by rule, capped at 10. So +4.0 moves 4.0 points, after the pool:

+25 bp: 20.2% → 22.4%

+50 bp 6.0 → 7.7; no change 70.9 → 67.6; cuts 3.0 → 2.3

ONE VERIFIED CLAIM, AS RECORDED · ANALYST 4, COMMUNICATIONS, RUN 1 OF 5

Claim The September statement stressed elevated inflation and a resolve to return to target on a timelier basis.

"Inflation remains elevated. Today's policy action will support a timelier return to the Committee's 2 percent goal. The Committee will deliver price stability."

Item **fed.statement.2026-09-16** · quote found by code, character for character · lean hawkish · 12 of 12 claims in this run verified

Live demo 4: what the analysts concluded

2 min

Open

- Analyst results view
- Analyst 1 (naive) card
- Analyst 4 (communications) card

Point at

- Analyst 1: cut 25 at 57%, numbers from memory flagged
- Analyst 4: the lean, grounded, leaked, watch-outs badges; one claim with its quote and the 'quote verified' mark
- Hover a badge for its definition

Say

- Every claim you see has a quote that code found character for character in the item named.
- The naive analyst still thinks the Fed is cutting. Same model, no evidence.

LIVE DEMO

Live demo 5: the scorecard and the drift table

2 min

Open

- Scorecard view
- The snapshot dropdown: pick 24 October, then 26 October
- Expand 'what was checked' under Shock direction
- Scroll to the over-time table

Point at

- The grid: red in the control columns, green in the judged columns
- The seven shocks with their deltas
- The over-time table: one row per snapshot, the marks never move

Say

- Analysts 1 and 2 are meant to fail. A red cell on the right would be a defect.
- This is the drift check. Two red cells have already pointed at defects and both are fixed.

LIVE DEMO

Live demo 6: the ledger and replay

1 min

Open

- Ledger view
- Press Verify ledger

Point at

- Every entry with its snapshot hash, spec hash and code version
- The replay result: chain intact, entries under the current spec recompute exactly
- Eight numbered amendments, each with its reason, each a new spec hash pinned in the tests

Say

- Nothing in this ledger can be edited or deleted without the chain breaking.
- Any entry can be recomputed from its snapshot. That is what auditable means here.

LIVE DEMO

What comes next

December is on file: entry #23 for the 9 December meeting, +25 bps 73.6%, no change 16.7%, specialists +4.0 bp. The 28 October call is scored the moment the post-meeting rate row lands.

1 Pre-registered weight sets, scored live

Write down three or four candidate weight sets now (for example market 0.6 / 0.5 / 0.7 with the rest redistributed). Score every set on every live decision with the same log score and Brier. The live set changes only by amendment, on that record. Weights stop being judgement without ever being fitted to the past.

2 Reasoning quality, graded by people

A grading form in the app: a user marks an analyst's answer from 1 to 5 on four questions (did it start from the base rate, argue both sides, say what it was unsure of, say what would prove it wrong). Twenty hand grades become the reference; a model grader is trusted only once it agrees with them. The grade becomes the tenth measured dimension and a gate on the lean.

3 An events agent for regime shifts

1987, 1998, 2008, 2020: the Fed cut on financial stress, not the data the models read. A watch-out agent that scans for bank failures, funding stress and market dislocations and raises a typed watch-out every analyst must answer.

4 A second model family

The same harness, the same snapshots, a different language model. If the ladder reads the same, the harness is doing the work, not the model.

5 Other central banks and geographies

ECB, Bank of England, MAS. The chain needs four things: a meeting calendar, a policy-rate series, the bank's own statements and minutes, and a futures or swap curve. The scorecard, the ledger and the analysts transfer unchanged; only the evidence sources change.

6 Better classical models

The probit's successor: features written down before fitting, a regularised ordinal model, a reliability check beside the log score, and the expanding window as the only judge. A regularised boosted model as one pre-registered candidate. About 130 usable decisions, 68% holds: small models, honestly tested, not big ones.

Three things to take away

- The harness is the product. Same model, same evidence, same day: naive it wobbles and leaks; harnessed it is stable, quoted and checked.
- Judge the process. Eight decisions a year cannot prove skill; a pre-registered scorecard measured on every snapshot can show discipline and catch drift.
- The ask: scrutinise the scorecard, challenge a quote, and tell us which other decisions deserve the same chain. The 28 October call is on file; the 9 December forecast is too.

Backup: the ladder, four snapshots, five runs each

Dimension (pass mark)	Analyst 1: the question only	Analyst 2: + the evidence	Analyst 3: + answer shape + verifier	Analysts 4-8: + one slice each
Repeatability (≤ 2 pp) largest standard deviation of any outcome probability across 5 identical runs	**3.4 / 4.5 / 4.9 / 4.1 pp, fail**	**3.3 / 2.3 / 4.8 / 3.4 pp, fail**	**0.6 / 0.5 / 0.5 / 0.8 bp of tilt, pass**	**0.5 / 0.8 / 1.2 / 0.8 bp of tilt, pass**
Quotes verified (≥ 95%) share of claims whose quote code found in the evidence	**0%, fail**	not measurable: no quotes	**100 / 100 / 100 / 99%, pass**	**100 / 97 / 100 / 100%, pass**
Numbers from memory (exactly 0) numbers in the answer that are not in the snapshot	**36 / 36 / 39 / 38, fail**	**0, pass**	**0, pass**	**0, pass**
What it concluded	Cut 25 bps at 57 to 67%: its memory of the Fed is stale	+25 bps at 50 to 57%: reads the evidence, overshoots the market	+7.0 / +6.6 / +6.4 / +5.8 bp hawkish	+4.4 / +4.1 / +4.4 / +4.0 bp hawkish; labour the lone non-hawk

Snapshots of 15, 24, 26 and 28 October 2026. Analysts 3 to 8 report repeatability in basis points of their direction call because they output a lean, not a distribution.