FedCast

A forecaster for the next Federal Reserve rate decision, built to be judged by its process rather than its luck

Tan Kah Ming · Fullerton · October 2026

About me

Tan Kah Ming

Senior AI, Data and Technology Consultant, seconded to Temasek Holdings, Digital Innovation (2023 to now)

Before that

- Data scientist, RSAF Data Office: authored the Air Force's AI strategy, set up its first digital unit
- Air engineering and weapons logistics officer, where the data-driven habit started (Weibull analysis, predictive maintenance)

Training

M.Eng Aeronautical Engineering, Imperial College London · Graduate Certificate in Predictive Analytics, NUS · CFA Level 3 candidate (February 2027)

Things I have built that models had to earn trust in

1

Investment analytics at Temasek

An AI-assisted analyst performance and attribution platform for 50+ investment professionals, with an LLM signal tracker for high-conviction ideas; signal cadence moved from monthly to daily, every recommendation logged in real time.

2

Valuation and execution

A machine-learning mark-to-market for private holdings, systematic and reproducible; and an AI-driven order management system on AWS processing about S$30m a day, replacing an email-based workflow.

3

Forecasting under a scoreboard

4th of 160+ teams (MIT, DoD, USAF) in the 2023 USAF AI Accelerator, predicting pilot flying difficulty from physiology; an ammunition recommender (quantile LightGBM) that cut wastage with no shortfall, Defence Manpower Committee Innovation Award 2022.

4

COVID-19, 2021

Built, in weeks, the system that engaged every home-recovery patient in Singapore, 10,000 people a day, on form.sg, Azure and Zendesk. The lesson that carries over: a system people rely on must say what it knows and what it does not.

What FedCast is, and why we built it

What it is

- A forecast of the next FOMC decision as probabilities over five outcomes, from one frozen, fingerprinted snapshot of evidence: futures, FRED series as first published, statements, minutes, transcript, projections, speeches, 32 years of decisions, prediction-market odds.
- Code computes the signals and pools four models with fixed weights: market-implied 0.6, the Fed's own Taylor rules 0.2, prediction-market odds 0.1, historical base rates 0.1.
- A language model is used only to read text: five specialist analysts, tools off, every claim quoted and verified by code, whose combined lean can tilt the number by at most 10 points and never write it.
- Human judgement enters only through typed, bounded, recorded channels: views, watch-outs, documents, odds.
- Every forecast goes into a hash-chained ledger that recomputes exactly from its snapshot.

Today, for 28 October 2026: no change 71%, +25 bps 23%

Why we built it

1 Accuracy is the wrong yardstick at this horizon

Eight decisions a year, and the futures market already calls about 88% of them a month out. A good year proves nothing. So the system is judged on a ten-dimension process scorecard fixed before any code existed, measured by code on every snapshot, and outcomes are logged but never used to tune.

2 A desk needs a number it can challenge

Every input is visible, every analyst claim carries a quote code has matched, every ledger entry replays. When the number moves, you can see which input moved it and by how much.

3 Language models are entering this work unchecked

Recent papers report 90%+ accuracy from models that read Fed text, on sixteen meetings, with no frozen evidence, no verification and no control. We built the control: the same model with no harness, with evidence only, with structure and verification, with specialisation, so the value of each technique is measured, not asserted.

4 The harness is the asset

The same chain can target the December meeting, the year-end rate or another central bank. The first live decision is scored this week, and the record starts there.