FedCast

A forecaster for the next Federal Reserve rate decision, built to be judged by its process rather than its luck

Tan Kah Ming · Fullerton · October 2026

About me

Tan Kah Ming

Senior AI, Data and Technology Consultant, seconded to Temasek Holdings, Digital Innovation (2023 to now)

Before that

- Data scientist, RSAF Data Office: authored the Air Force's AI strategy, set up its first digital unit
- Air engineering and weapons logistics officer, where the data-driven habit started (Weibull analysis, predictive maintenance)

Training

M.Eng Aeronautical Engineering, Imperial College London · Graduate Certificate in Predictive Analytics, NUS · CFA Level 3 candidate (February 2027)

Things I have built that models had to earn trust in

1

Investment analytics at Temasek

An AI-assisted analyst performance and attribution platform for 50+ investment professionals, with an LLM signal tracker for high-conviction ideas; signal cadence moved from monthly to daily, every recommendation logged in real time.

2

Valuation and execution

A machine-learning mark-to-market for private holdings, systematic and reproducible; and an AI-driven order management system on AWS processing about S$30m a day, replacing an email-based workflow.

3

Forecasting under a scoreboard

4th of 160+ teams (MIT, DoD, USAF) in the 2023 USAF AI Accelerator, predicting pilot flying difficulty from physiology; an ammunition recommender (quantile LightGBM) that cut wastage with no shortfall, Defence Manpower Committee Innovation Award 2022.

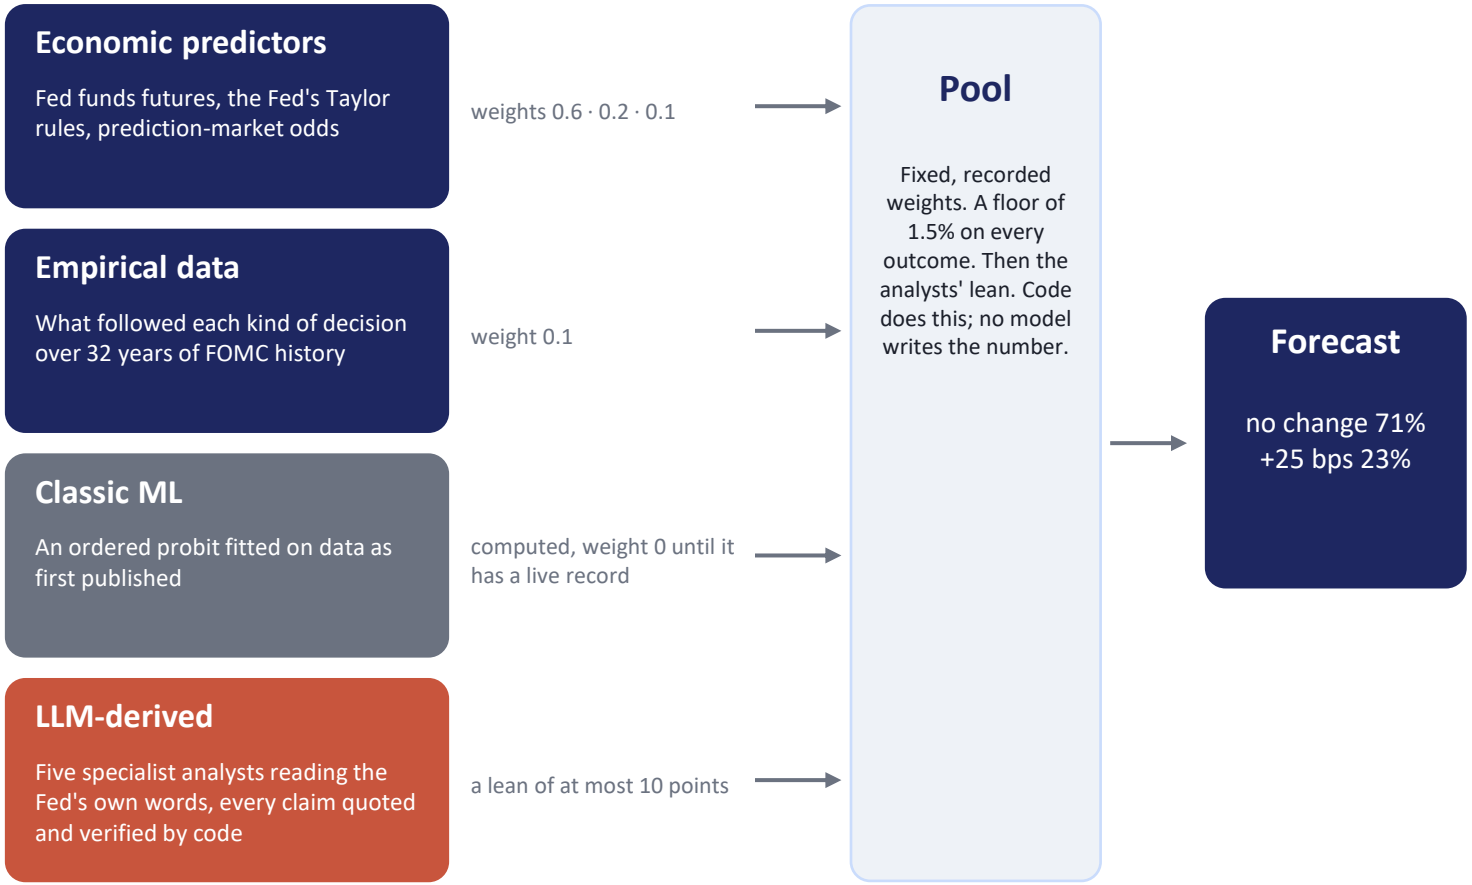
4

COVID-19, 2021

Built, in weeks, the system that engaged every home-recovery patient in Singapore, 10,000 people a day, on form.sg, Azure and Zendesk. The lesson that carries over: a system people rely on must say what it knows and what it does not.

What FedCast is, and why we built it

A forecast of the next FOMC decision as probabilities over five outcomes, from one frozen snapshot of evidence. Today, for 28 October 2026: no change 71%, +25 bps 23%.



Repeatability is enforced, then measured

One frozen, fingerprinted snapshot per forecast; code pools the models; each analyst is run five times on identical inputs with every tool off and every quote checked by code. A ten-dimension scorecard, fixed before any code existed, is measured on every snapshot, so drift shows as a red cell.

Human judgement has a typed channel

None is active today. The channel exists because a desk's judgement belongs in the number, but only where it can be seen, bounded and scored afterwards: a view tilts the number by at most 10 points and expires; a watch-out must be answered by every analyst; a document or an odds entry becomes evidence.

All evidence is kept, so the path can be read back

Each pull is frozen and hash-chained to the forecast it produced. As new information arrives, the number and each model's contribution move, and the ledger shows exactly which input moved it: 62/33 on 15 October, 65/30 on 24 October, 71/23 on 26 October after the market odds were refreshed.